

Class07 Machine Learning 1

Zixuan Zeng (A16142927)

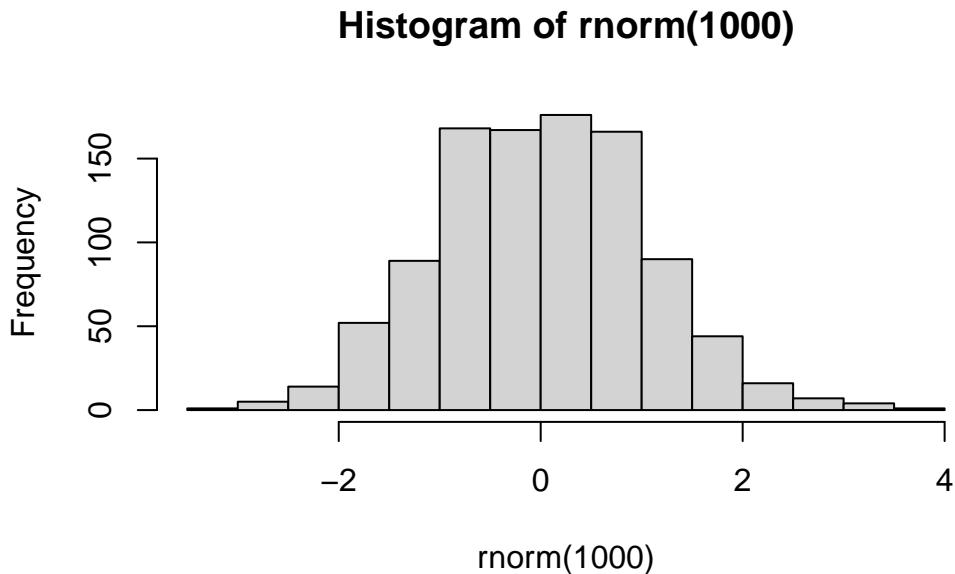
Table of contents

Hierarchical Clustering	6
PCA	8
PCA to the rescue	13
RNA-seq data exploration	18

Today we will begin our exploration of some “classical” machine learning approaches. We will start with clustering:

Let’s first make up some data to cluster where we know what the answer should be.

```
hist( rnorm(1000) )
```



```

x <- c(rnorm(30,mean=-3), rnorm(30,mean=3))
y <- rev(x)

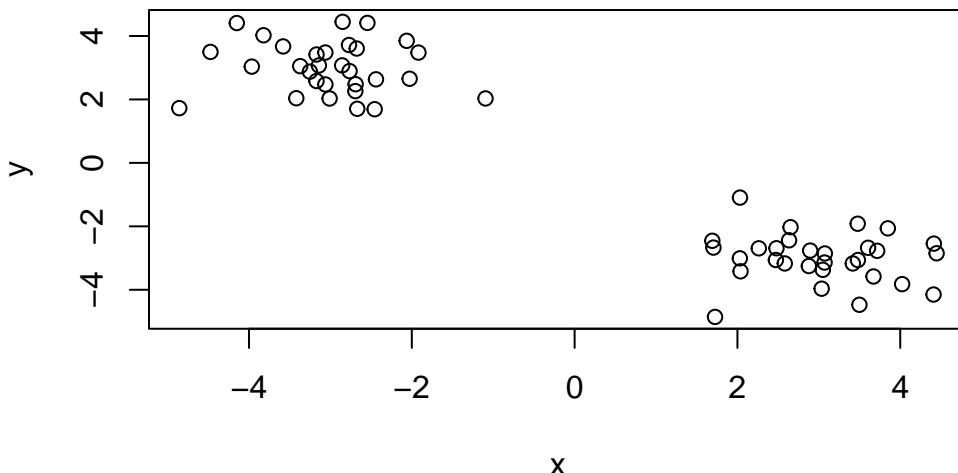
x <- cbind(x,y)
head(x)

```

	x	y
[1,]	-2.545408	4.412533
[2,]	-3.581017	3.671296
[3,]	-2.855466	3.073373
[4,]	-3.821666	4.023172
[5,]	-3.418836	2.038392
[6,]	-2.764325	2.893790

plot x

```
plot(x)
```



The main function in “base” R for K-means clustering is called `kmeans()`.

```
k <- kmeans(x, centers=2)  
k
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```

          x           y
1 -3.000652  3.010232
2  3.010232 -3.000652

```

Clustering vector:

Within cluster sum of squares by cluster:

```
[1] 36.76886 36.76886
```

(between SS / total SS = 93.6 %)

Available components:

```
[1] "cluster"      "centers"       "totss"        "withinss"      "tot.withinss"  
[6] "betweenss"    "size"          "iter"         "ifault"
```

Q. How big are the clusters (i.e. their size)?

k\$size

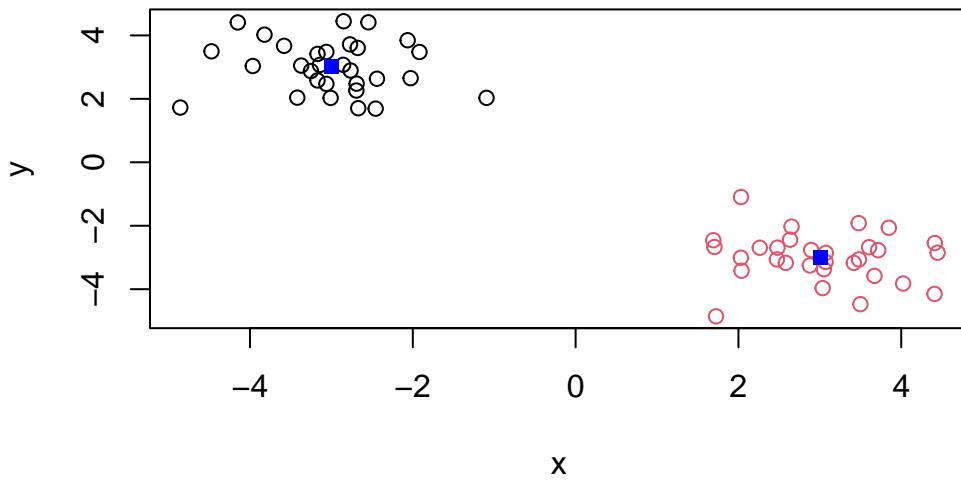
[1] 30 30

Q. What clusters do my data points reside in?

k\$cluster

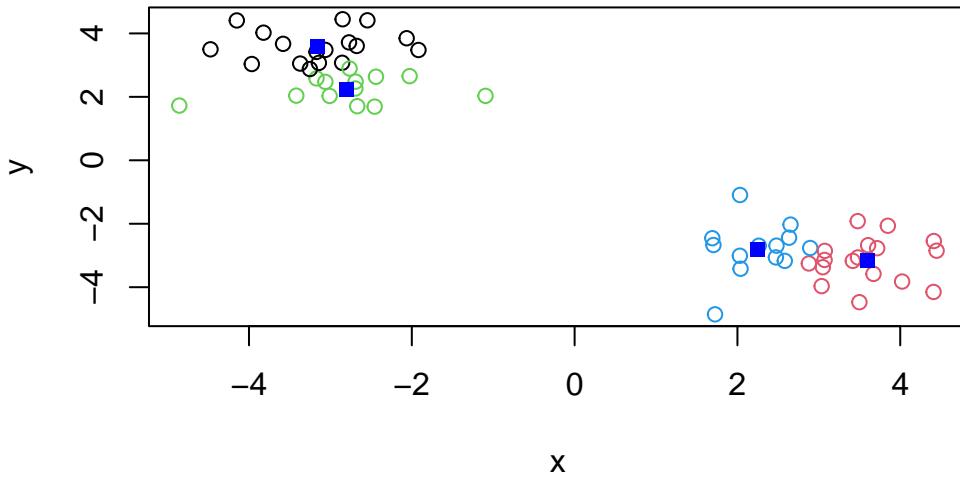
Q. Make a plot of our data colored by cluster assignment(i.e. make a result figure)

```
plot(x,col=k$cluster)
points(k$centers,col="blue",pch=15)
```



Q. Cluster with k-means into 4 clustrs and plot your results as above.

```
k4 <- kmeans(x, centers=4)
plot(x,col=k4$cluster)
points(k4$centers,col="blue",pch=15)
```



Q. Run kmeans with center (i.e. values of k) equal 1 to 6

```
k1 <- kmeans(x, centers=1)$tot.withinss
k2 <- kmeans(x, centers=2)$tot.withinss
k3 <- kmeans(x, centers=3)$tot.withinss
k4 <- kmeans(x, centers=4)$tot.withinss
k5 <- kmeans(x, centers=5)$tot.withinss
k6 <- kmeans(x, centers=6)$tot.withinss

ans <- c(k1,k2,k3,k4,k5,k6)
```

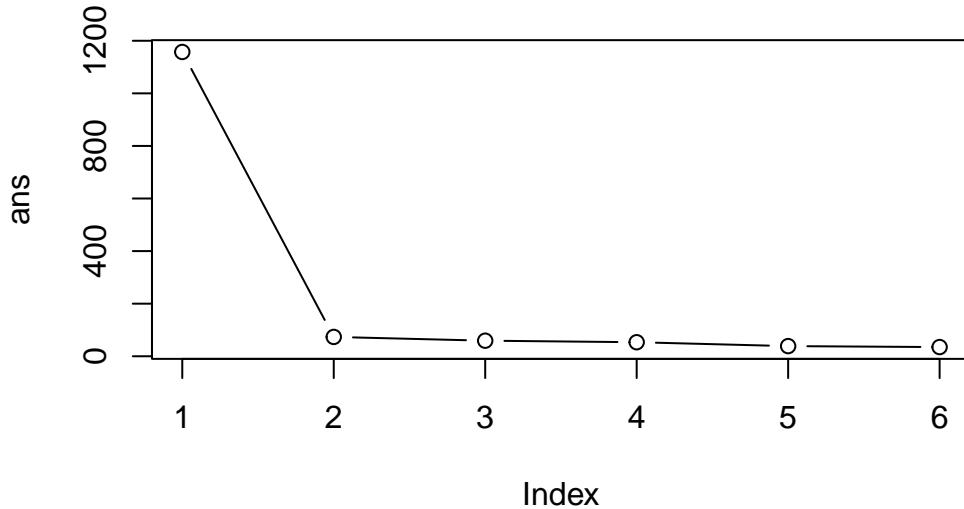
Or use a **for** loop

```
ans <- NULL
for(i in 1:6){
  km <- kmeans(x, centers=i)
  ans[i] <- km$tot.withinss
}
ans
```

```
[1] 1157.45952   73.53773   59.17998   53.34115   38.60985   34.91535
```

Make a “screen-plot”

```
plot(ans,typ="b")
```



Hierarchical Clustering

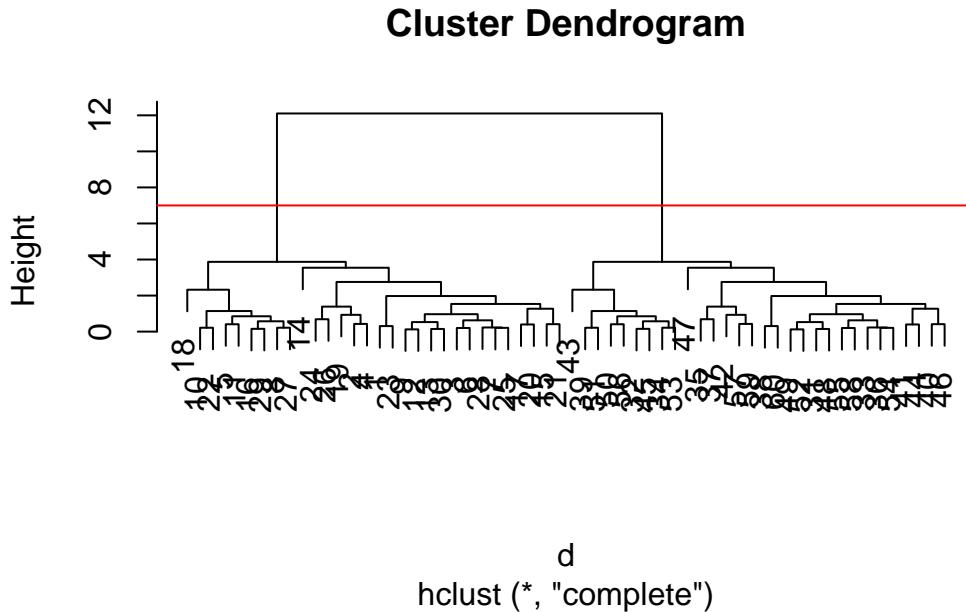
The main function in “base” R for this is called `hclust()`.

```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:
`hclust(d = d)`

Cluster method : complete
Distance : euclidean
Number of objects: 60

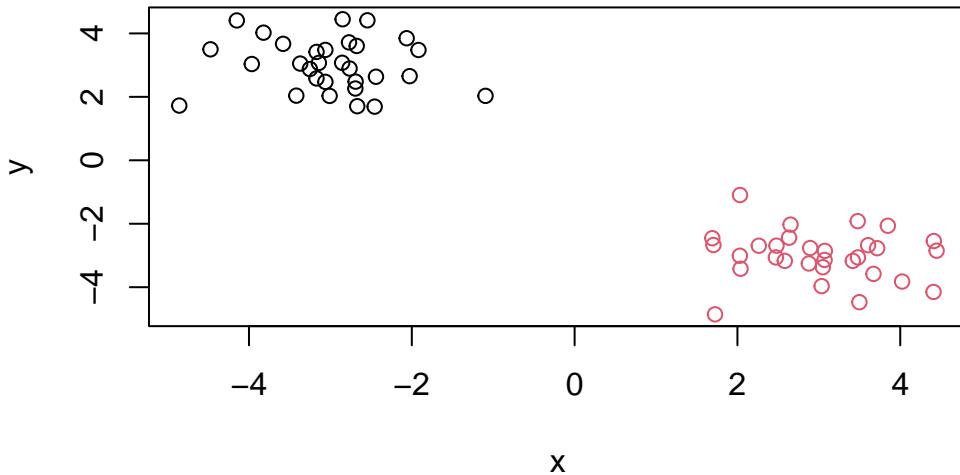
```
plot(hc)
abline(h=7, col="red")
```



To obtain clusters from our `hclust` result object `hc`, we cut the tree to yield different sub branches. For this we use the `cutree()` function.

```
grps <- cutree(hc,h=7)  
grps
```

```
plot(x,col=grps)
```



PCA

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. 17 rows and 5 column dim(x)

Q2. I like the first approach as the code feels more straight-forward to me. Running it more than one time could remove the information we want.

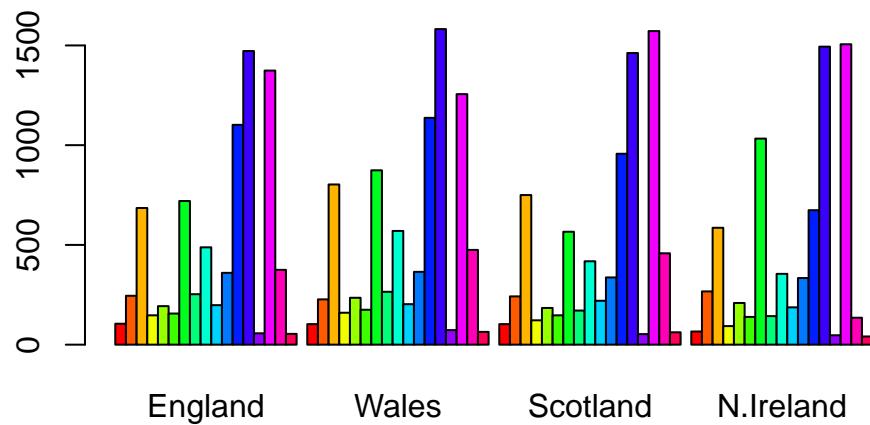
Q3. beside = FALSE

```
rownames(x) <- x[,1]
x <- x[,-1]
head(x)
```

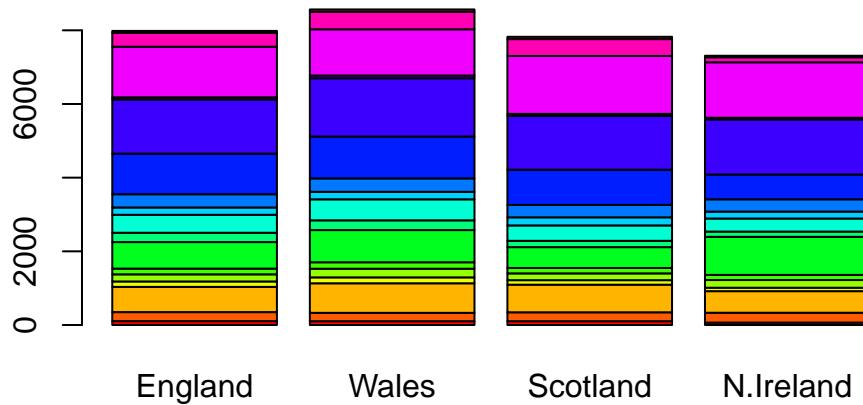
	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93

Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```

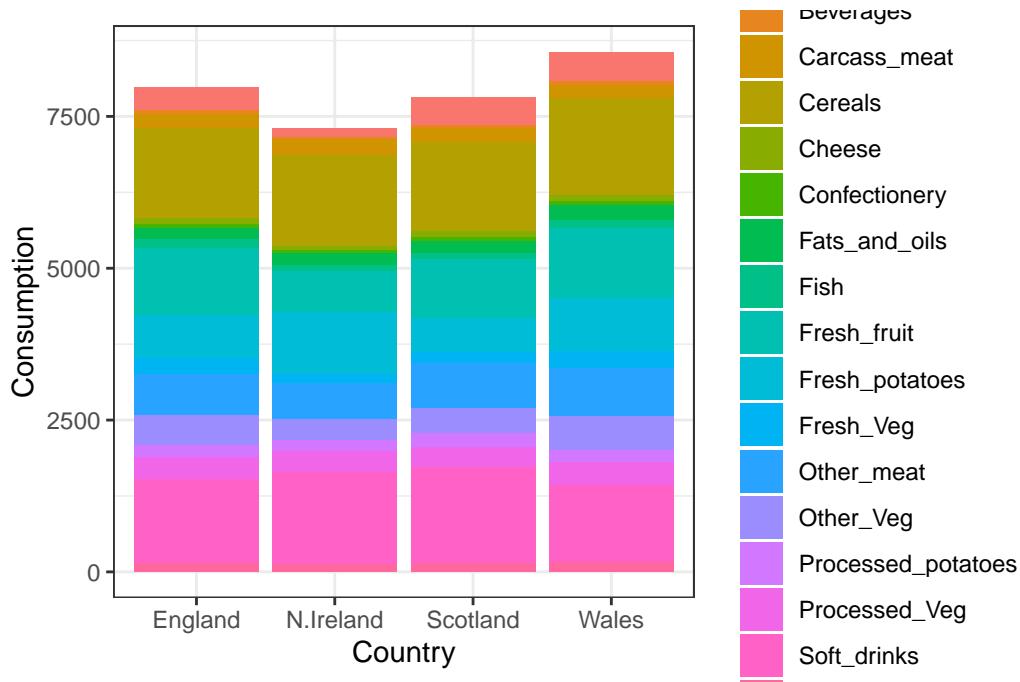


```
barplot(as.matrix(x), beside=FALSE, col=rainbow(nrow(x)))
```



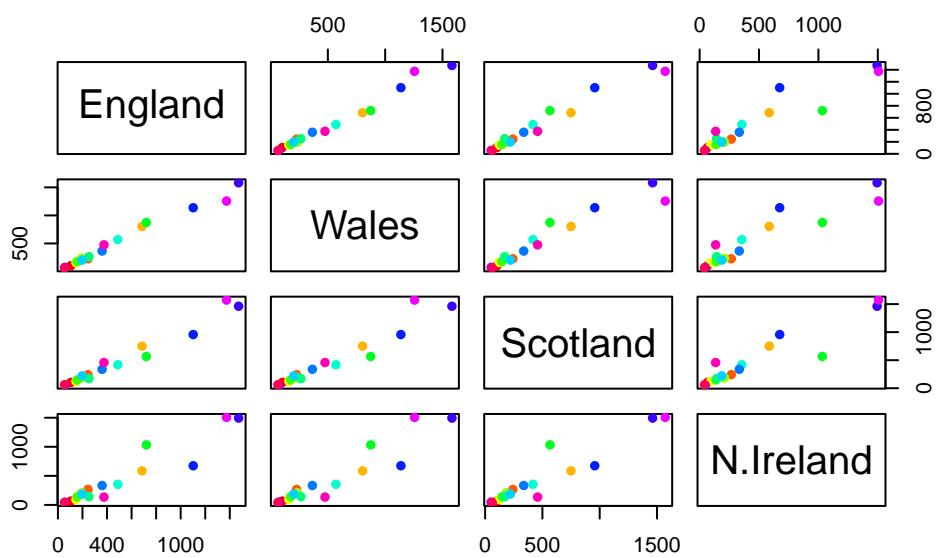
Q4. Delete the position = “dodge” argument

```
library(tidyr)
library(ggplot2)
x_long <- x |>
  tibble::rownames_to_column("Food") |>
  pivot_longer(cols = -Food,
                names_to = "Country",
                values_to = "Consumption")
ggplot(x_long) +
  aes(x = Country, y = Consumption, fill = Food) +
  geom_col() +
  theme_bw()
```

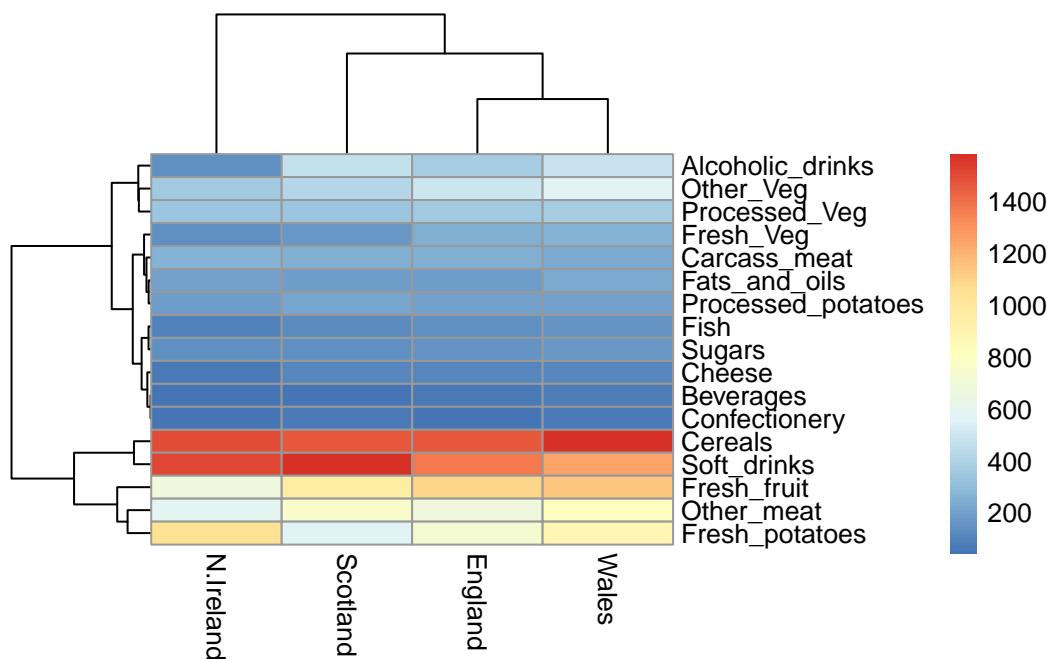


Q5. Plots above and below each diagonal countries are replicates. Each plot represents comparison between the two countries vertical to the plot. Each dot represents a food item. The position of the dot represents the consumption amount in each country. The dot would be on the diagonal of the plot if two countries consume about the same amount of that particular food.

```
pairs(x, col=rainbow(nrow(x)), pch=16)
```



```
library(pheatmap)
pheatmap( as.matrix(x) )
```



Q6. Scotland, England, and Wales are clustered together in contrast to N.Ireland. This result indicates that N.Ireland has a different food consumption pattern than the other three countries. Main differences are a little hard to spot since the color is not that different.

PCA to the rescue

The main function in “base” R for PCA is called `prcomp()`.

As we want to do a PCA on the food data for the different countries, we will want the foods in the columns.

```
PCA <- prcomp(t(x))
summary(PCA)
```

Importance of components:

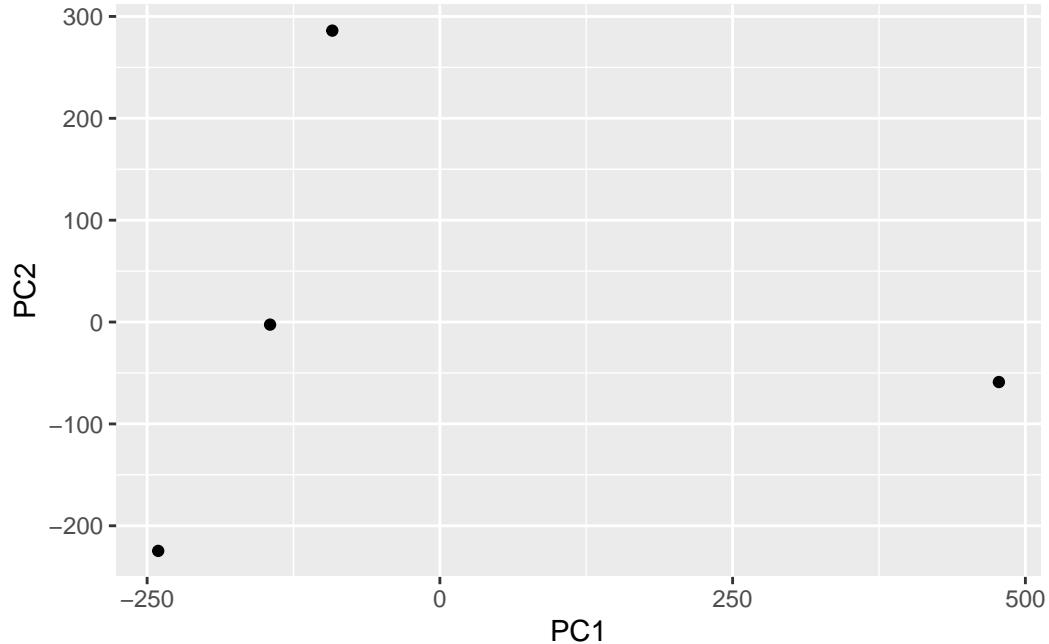
	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	2.921e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Our result object is called `PCA` and it has a `$x` component that we will look at first.

```
PCA$x
```

	PC1	PC2	PC3	PC4
England	-144.99315	-2.532999	105.768945	-9.152022e-15
Wales	-240.52915	-224.646925	-56.475555	5.560040e-13
Scotland	-91.86934	286.081786	-44.415495	-6.638419e-13
N.Ireland	477.39164	-58.901862	-4.877895	1.329771e-13

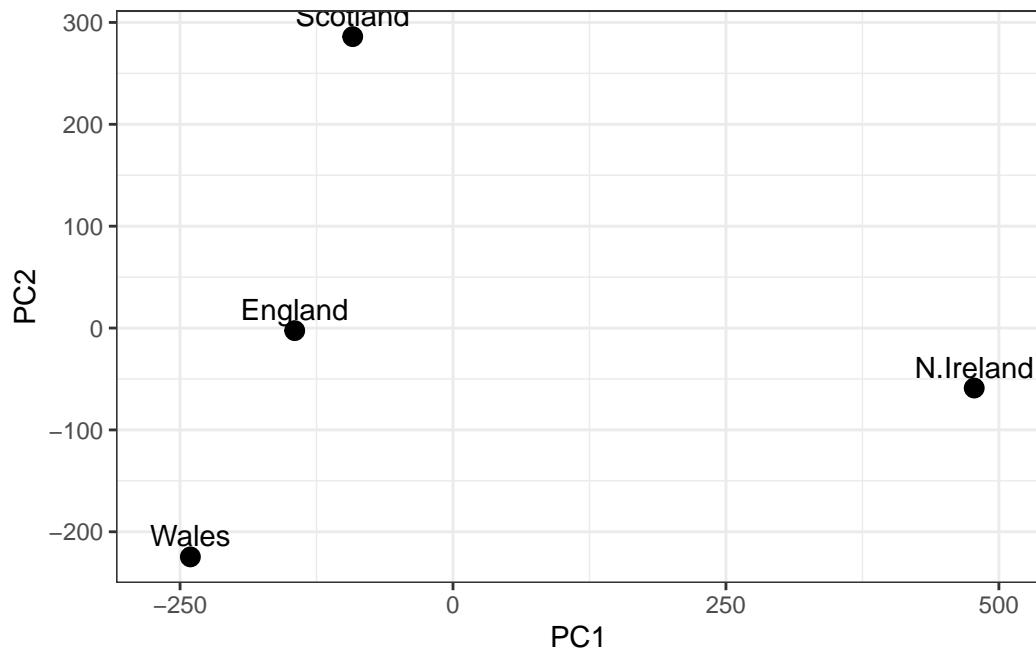
```
ggplot(PCA$x) + aes(PC1,PC2) + geom_point()
```



Q7. The finished code is below.

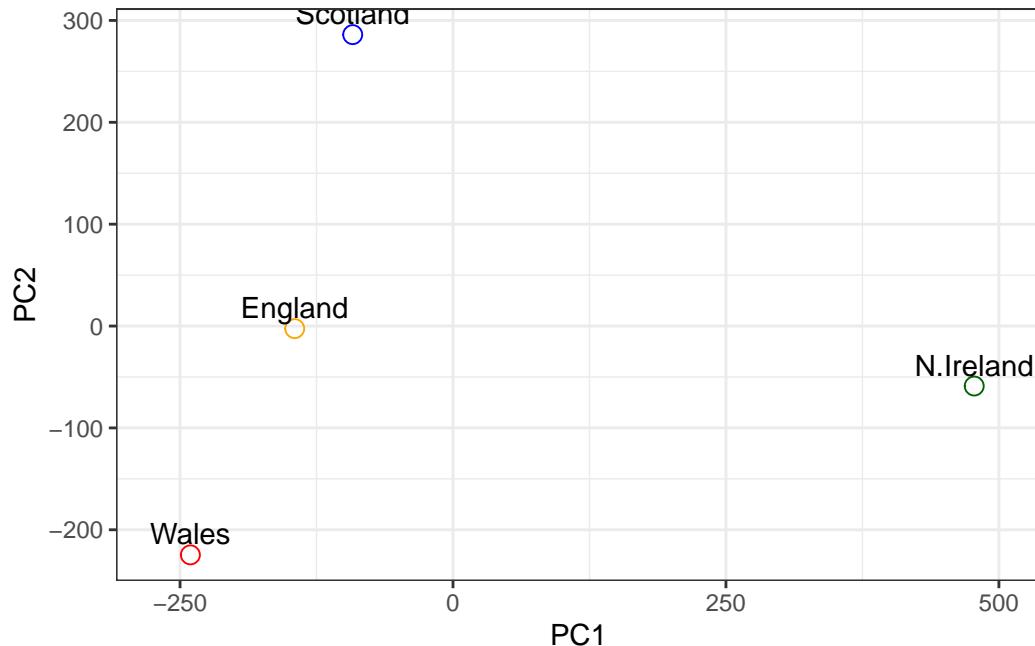
```
# Create a data frame for plotting
df <- as.data.frame(PCA$x)
df$Country <- rownames(df)

# Plot PC1 vs PC2 with ggplot
ggplot(PCA$x) +
  aes(x = PC1, y = PC2, label = rownames(PCA$x)) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5) +
  xlim(-270, 500) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()
```



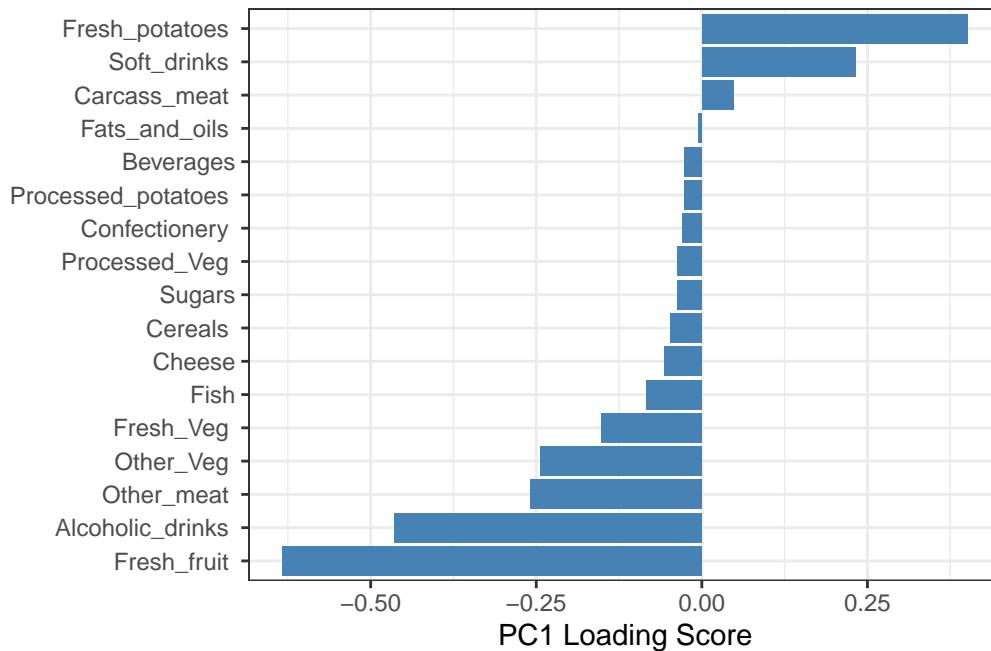
Q8.

```
ggplot(PCA$x) +  
  aes(x = PC1, y = PC2, label = rownames(PCA$x)) +  
  geom_point(size = 3, color=c('orange', 'red', 'blue', 'darkgreen'), pch = 1) +  
  geom_text(vjust = -0.5) +  
  xlim(-270, 500) +  
  xlab("PC1") +  
  ylab("PC2") +  
  theme_bw()
```



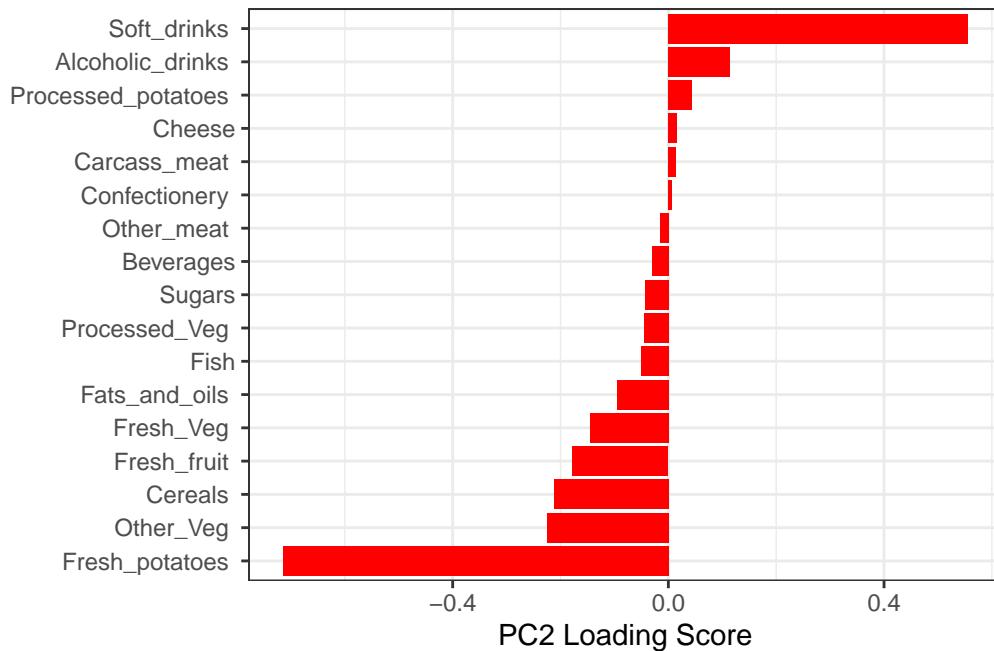
Another major result out of PCA is the so-called “variable loadings” or `$rotation` that tells us how the original variables (foods) contribute to PCs.

```
ggplot(PCA$rotation) +
  aes(x = PC1,
      y = reorder(rownames(PCA$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



Q9. Soft drinks and alcoholic drinks feature prominently in PC2. PC2 mainly tells us that consumption of soft and alcoholic drinks separates Scotland from the cluster that contains the rest of three countries.

```
ggplot(PCA$rotation) +
  aes(x = PC2,
      y = reorder(rownames(PCA$rotation), PC2)) +
  geom_col(fill = "red") +
  xlab("PC2 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))
```



RNA-seq data exploration

Q10. 100 genes and 10 samples

```
url2 <- "https://tinyurl.com/expression-CSV"
rna.data <- read.csv(url2, row.names=1)
head(rna.data)
```

	wt1	wt2	wt3	wt4	wt5	ko1	ko2	ko3	ko4	ko5
gene1	439	458	408	429	420	90	88	86	90	93
gene2	219	200	204	210	187	427	423	434	433	426
gene3	1006	989	1030	1017	973	252	237	238	226	210
gene4	783	792	829	856	760	849	856	835	885	894
gene5	181	249	204	244	225	277	305	272	270	279
gene6	460	502	491	491	493	612	594	577	618	638

```
dim(rna.data)
```

```
[1] 100 10
```

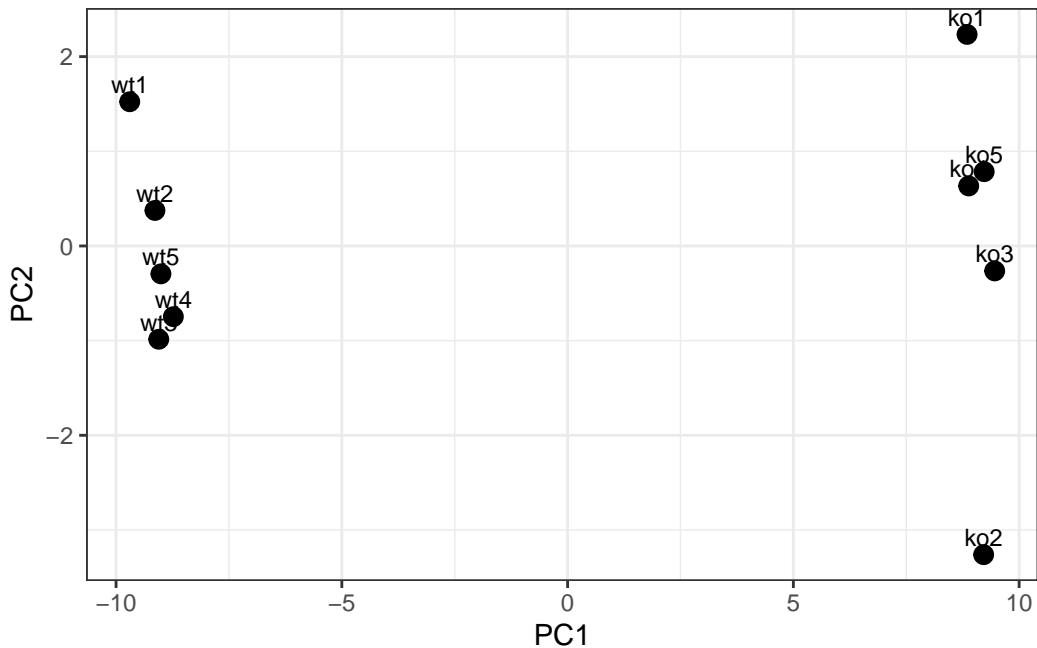
```

## Again we have to take the transpose of our data
pca <- prcomp(t(rna.data), scale=TRUE)

# Create data frame for plotting
df <- as.data.frame(pca$x)
df$Sample <- rownames(df)

## Plot with ggplot
ggplot(df) +
  aes(x = PC1, y = PC2, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, size = 3) +
  xlab("PC1") +
  ylab("PC2") +
  theme_bw()

```



```
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4	PC5	PC6	PC7
Standard deviation	9.6237	1.5198	1.05787	1.05203	0.88062	0.82545	0.80111
Proportion of Variance	0.9262	0.0231	0.01119	0.01107	0.00775	0.00681	0.00642

```

Cumulative Proportion  0.9262 0.9493 0.96045 0.97152 0.97928 0.98609 0.99251
                           PC8      PC9      PC10
Standard deviation     0.62065 0.60342 3.3e-15
Proportion of Variance 0.00385 0.00364 0.0e+00
Cumulative Proportion  0.99636 1.00000 1.0e+00

```

```

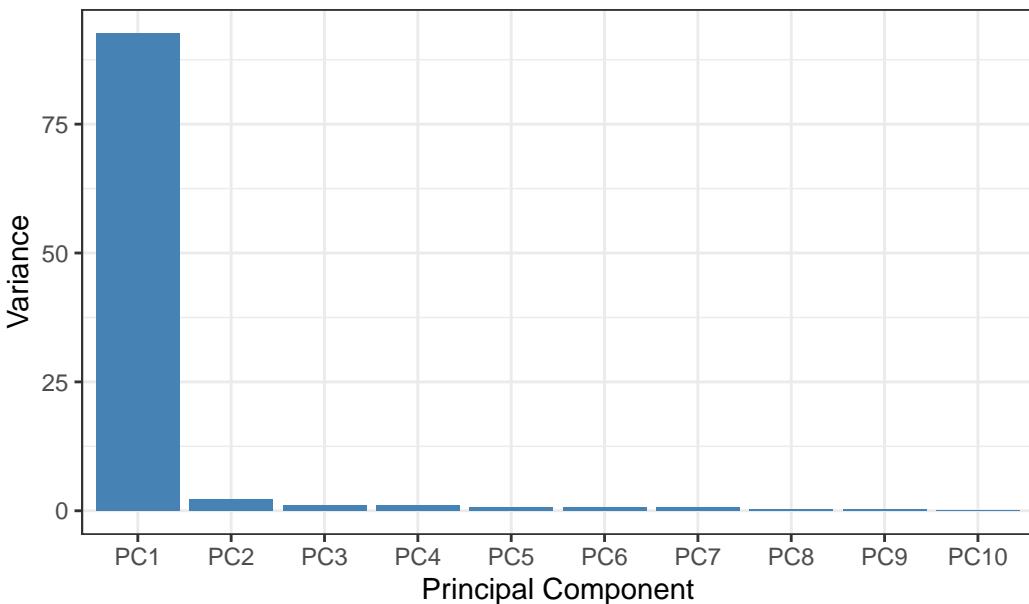
# Calculate variance explained
pca.var <- pca$sdev^2
pca.var.per <- round(pca.var/sum(pca.var)*100, 1)

# Create scree plot data
scree_df <- data.frame(
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),
  Variance = pca.var[1:10]
)

ggplot(scree_df) +
  aes(x = PC, y = Variance) +
  geom_col(fill = "steelblue") +
  ggtitle("Quick scree plot") +
  xlab("Principal Component") +
  ylab("Variance") +
  theme_bw()

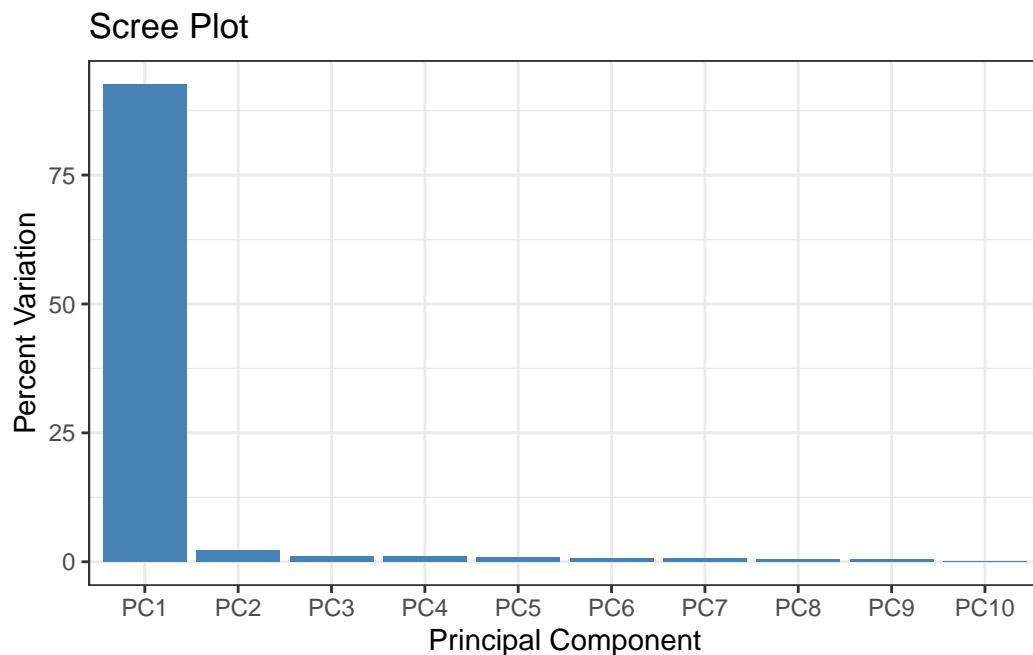
```

Quick scree plot



```
pca.var.per  
[1] 92.6 2.3 1.1 1.1 0.8 0.7 0.6 0.4 0.4 0.0
```

```
# Create percent variance scree plot  
scree_pct_df <- data.frame(  
  PC = factor(paste0("PC", 1:10), levels = paste0("PC", 1:10)),  
  PercentVariation = pca.var.per[1:10]  
)  
  
ggplot(scree_pct_df) +  
  aes(x = PC, y = PercentVariation) +  
  geom_col(fill = "steelblue") +  
  ggtitle("Scree Plot") +  
  xlab("Principal Component") +  
  ylab("Percent Variation") +  
  theme_bw()
```



```
## A vector of colors for wt and ko samples  
colvec <- colnames(rna.data)  
colvec[grep("wt", colvec)] <- "red"
```

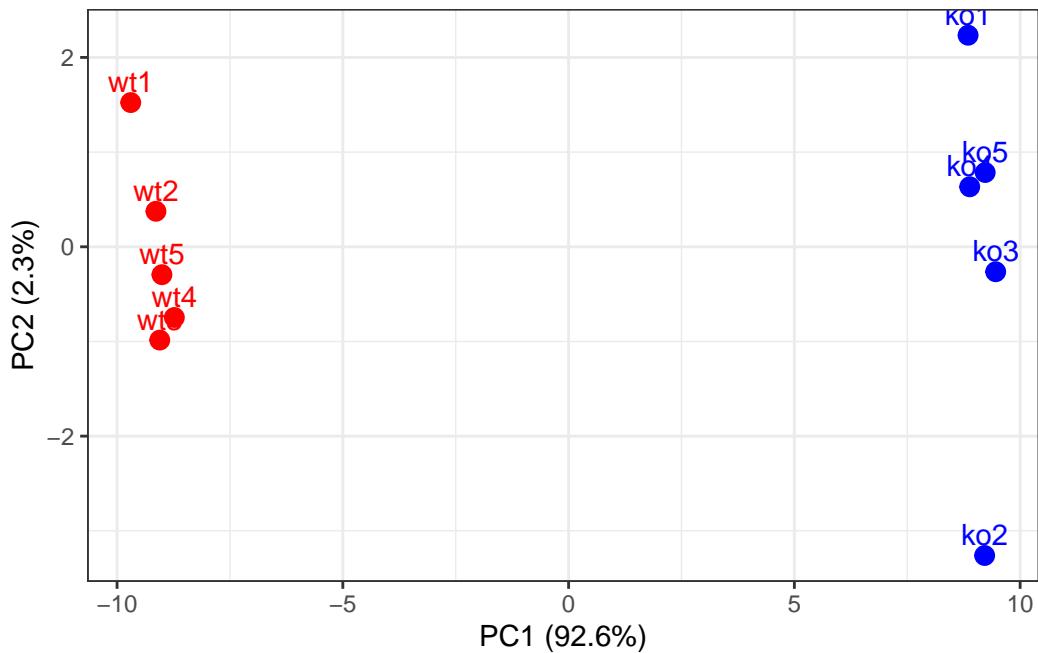
```

colvec[grep("ko", colvec)] <- "blue"

# Add condition to data frame
df$condition <- substr(df$Sample, 1, 2)
df$color <- colvec

ggplot(df) +
  aes(x = PC1, y = PC2, color = color, label = Sample) +
  geom_point(size = 3) +
  geom_text(vjust = -0.5, hjust = 0.5, show.legend = FALSE) +
  scale_color_identity() +
  xlab(paste0("PC1 (", pca.var.per[1], "%)")) +
  ylab(paste0("PC2 (", pca.var.per[2], "%)")) +
  theme_bw()

```



```

loading_scores <- pca$rotation[,1]

## Find the top 10 measurements (genes) that contribute
## most to PC1 in either direction (+ or -)
gene_scores <- abs(loading_scores)
gene_score_ranked <- sort(gene_scores, decreasing=TRUE)

```

```

## show the names of the top 10 genes
top_10_genes <- names(gene_score_ranked[1:10])
top_10_genes

[1] "gene100" "gene66"   "gene45"   "gene68"   "gene98"   "gene60"   "gene21"
[8] "gene56"   "gene10"   "gene90"

ggplot(pca$rotation) +
  aes(x = PC1,
      y = reorder(rownames(pca$rotation), PC1)) +
  geom_col(fill = "steelblue") +
  xlab("PC1 Loading Score") +
  ylab("") +
  theme_bw() +
  theme(axis.text.y = element_text(size = 9))

```

